

END-OF-COURSE PROJECT • MILESTONE 1

Project Problem Statement

Submit by the end of Week 3 [23 August 23:59 hrs] • ~1 page • individual • formative (not graded) — you'll get written feedback. This seeds your End-of-Course Project (it maps to Project Rubric criteria 1 & 2).

This submission is NOT graded. However, on time non-submission counts towards incomplete End of course project.

Name: Yang Ruijia Section (A / B / C): A Date: 23 August 2026

1 • Working title

A short, specific name for your project.

NoteTriage — splits client minutes into fact, assumption, and open question.

2 • The problem, and why it matters

Who has this problem, and why is it worth solving? 2–3 sentences. (Rubric 1)

After a client discovery meeting, product and implementation staff have to turn messy notes into a brief the delivery team can trust. The failure is specific: guesses and unanswered questions get written as if they were client facts, and a 6–8 week IT pilot then plans around an API, a dataset, or a policy that was never confirmed. Closest tools — ChatGPT, Notion AI, Fireflies — summarise fluently; they do not force a fact / assumption / to-confirm split, and they do not quote the line they relied on. One mixed-up “fact” is enough to waste a sprint.

3 • Who it's for, and the domain

Your user / audience and the industry (e.g. banking, healthcare, retail, your employer).

Primary user: an implementation / product associate at a B2B IT services firm (enterprise and public-sector digitalisation). Persona: Mei, a first-year implementation associate, writing tomorrow's business-understanding brief from a one-hour discovery transcript at 10pm; she cannot go back to the client that night, and she must not tell the delivery lead a guess as a fact.

4 • Why AI — and which kind? Remember, from class 1 workshops, no one size fits all

Is AI the right tool here (vs a rule or lookup)? If yes, which type — narrow ML, a foundation model, RAG, an agent? (ILO 1)

A keyword rule (maybe, should, TBD, “I think”) catches some assumptions and misses the dangerous ones — statements said as facts that are actually unverified (“they already have SSO”). Lookup tables cannot classify unseen phrasing. So AI is warranted, but only as a three-way sentence classifier: a prompted foundation model with JSON output. Non-AI baseline: those keyword/regex rules, scored on the same labels. Not RAG in v1 (one document, no corpus). Not an agent (no tools, no loop). Climb later only if retrieval over past briefs earns it.

5 • Proposed approach — build vs buy or both!

What will you build (low-code and/or code)? Which layers will you own vs rent, and why? (Rubric 2, ILO 3)

Code, not low-code: a Python notebook. I tried the cheaper route first — pasting notes into ChatGPT and Notion AI. They summarised well but would not lock a three-label schema or refuse a “fact” without a quote; that is what pushed me to code. Rent the model via the course OpenRouter key (commodity generation). Own the schema, the quoting rule, the keyword baseline, the labelled eval set, and the guardrail “no quote → cannot be a fact.” No product UI for this course.

6 • Data – this packs the intelligence !

What data do you need, where from, and any access / privacy / residency constraints?

A labelled set of ~50 short discovery notes (about 400–600 sentences) in the IT-delivery domain (accounts, VPN, OA, integrations). Gold labels are frozen in a file before any model run: fact / assumption / open_question, plus the source span. I will write ~15 notes from intern-style scenes I already know, then generate the rest with a committed script and hand-correct labels. No live client transcripts; no names, emails, or employer files (PDPA). The repo will ship the generator, a sample, and the labels.

7 • Success metric & how you'll evaluate it

One metric that ties to real value, and one check (an eval) you could actually run. (evals thread)

Primary metric: assumption-as-fact rate on a 15-note holdout I labelled before seeing model outputs. Target: below 10%, versus the keyword baseline (value = fewer false commitments into the delivery brief).

Eval I can actually run: sentence-level macro-F1 on the three classes, plus a machine check that every “fact” row contains a verbatim quote from the note. The system may abstain; I will report whether abstentions fall on the cases it would have got wrong.

8 • Risks, limitations & responsible use

What could go wrong (incl. silent failure), and who might it affect?

Silent failure: labelling an assumption as a fact, which a delivery lead then treats as a client commitment. Mitigation: facts require a quote; low confidence → to-confirm; a human must accept the table before it leaves the notebook. Privacy: redact person and organisation names before the API call; course key only; no real client data (PDPA). Synthetic data has a ceiling — I will spot-check 10 generated notes against my own anonymised intern briefs and state what the set cannot tell us. This is an assistive classifier, not an autonomous commit (IMDA Model AI Governance Framework).

9 • Smallest first version (optional)

The tiniest end-to-end slice you could build first.

One pasted meeting note in → one foundation-model call with a fixed JSON schema out → a three-column table (fact / assumption / to-confirm), each row with a quote, shown next to the keyword-baseline table. That is the whole first slice.

How this is used

This milestone is not graded — it’s to give you fast, written feedback while there’s still time to change direction. A strong problem statement makes criteria 1 (problem & significance) and 2 (business & technical trade-offs) and helps you start on a strong footing. Bring questions during class breakouts or email the course account or WA text me

INSTRUCTOR FEEDBACK

For marker use — formative comments on the problem & significance (criterion 1) and the business/technical trade-offs (criterion 2), plus one thing to strengthen before the final project.